

Regularization
Gradient Descent with Momentum
RMSprop
Adam

13th December 2021

Overview

- 1 Regularization
- 2 Mini-batch gradient descent
- 3 Gradient descent with momentum
- 4 RMSprop
- 5 Adam

Regularization

What is regularization and why do we need it?




Figure: Bias-Variance ¹

¹<https://www.coursera.org/learn/deep-neural-network>

Regularization - Logistic Regression

$$z^{(i)} = w_0 + w_1 x_1^{(i)} + w_2 x_2^{(i)} + \dots + w_d x_d^{(i)}$$
$$a^{(i)} = \sigma(z^{(i)})$$

L2 regularization:

$$\min_{\mathbf{w}} L2(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N L^{(i)}(a^{(i)}, y^{(i)}) + \frac{\lambda}{2N} \|\mathbf{w}\|_2^2$$

where $\|\mathbf{w}\|_2^2 = \sum_{j=1}^d w_j^2$.

L1 regularization:

$$\min_{\mathbf{w}} L1(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N L^{(i)}(a^{(i)}, y^{(i)}) + \frac{\lambda}{2N} \|\mathbf{w}\|_1$$

where $\|\mathbf{w}\|_1 = \sum_{j=1}^d |w_j|$.

Note: We can omit regularization on w_0 .

Regularization - Neural Network

$$a^{(i)} = f^{[K]}(\mathbf{W}^{[K]}(\dots(f^{[2]}(\mathbf{W}^{[2]}(f^{[1]}(\mathbf{W}^{[1]}\mathbf{x}^{(i)}))))))$$

L2 regularization:

$$\min_{\mathbf{w}} L2(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N L^{(i)}(a^{(i)}, y^{(i)}) + \frac{\lambda}{2N} \sum_{k=1}^K \|\mathbf{W}^{[k]}\|_F^2$$

where $\|\mathbf{W}^{[k]}\|_2^2 = \sum_{i=1}^{n^{[k]}} \sum_{j=1}^{n^{[k-1]}} (w_{ij}^{[k]})^2$ is the Frobenius norm of the weight matrix in layer k .

L1 regularization:

$$\min_{\mathbf{w}} L1(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N L^{(i)}(a^{(i)}, y^{(i)}) + \frac{\lambda}{2N} \sum_{k=1}^K \|\mathbf{W}^{[k]}\|_1$$

where $\|\mathbf{W}^{[k]}\|_1 = \sum_{i=1}^{n^{[k]}} \sum_{j=1}^{n^{[k-1]}} |w_{ij}^{[k]}|$.

Note: We can omit regularization on $w_{i0}^{[k]}$.

Regularization - Neural Network

$\nabla L(\mathbf{W}^{[k]})$ is the gradient of the loss function without regularization obtained from backpropagation.

The gradient of the loss function with L2 regularization can be computed to update the weight matrix of layer k as:

$$\begin{aligned}\nabla L_2(\mathbf{W}^{[k]}) &\leftarrow \nabla L(\mathbf{W}^{[k]}) + \frac{\lambda}{N} \mathbf{W}^{[k]} \\ \mathbf{W}^{[k]} &\leftarrow \mathbf{W}^{[k]} - \gamma \nabla L_2(\mathbf{W}^{[k]}) \\ \mathbf{W}^{[k]} &\leftarrow \mathbf{W}^{[k]} - \gamma \left(\nabla L(\mathbf{W}^{[k]}) + \frac{\lambda}{N} \mathbf{W}^{[k]} \right) \\ \mathbf{W}^{[k]} &\leftarrow \left(1 - \frac{\gamma \lambda}{N} \right) \mathbf{W}^{[k]} - \gamma \nabla L(\mathbf{W}^{[k]})\end{aligned}$$

Gradient descent for L -layer neural network

Result: weights $\mathbf{W}^{[k]}$ for all layers.

for $k \leftarrow 1$ **to** L **do**

$\mathbf{W}^{[k]} \leftarrow \text{random}()$ *[Random initialization]*

end

for $t \leftarrow 1$ **to** max_iterations **do**

for $k \leftarrow 1$ **to** L **do**

$\mathbf{A}^{[k]} \leftarrow f(\mathbf{Z}^{[k]}) \leftarrow f(\mathbf{W}^{[k]} \cdot \mathbf{A}^{[k-1]})$ *[Forward propagation]*

end

for $k \leftarrow L$ **to** 1 **do**

$\nabla L(\mathbf{Z}^{[k]}) \leftarrow \nabla L(\mathbf{A}^{[k]}) \circ f'(\mathbf{Z}^{[k]})$

$\nabla L(\mathbf{W}^{[k]}) \leftarrow \frac{1}{N} \nabla L(\mathbf{Z}^{[k]}) \cdot \mathbf{A}^{[k-1]T}$ *[Backpropagation]*

$\nabla L(\mathbf{A}^{[k-1]}) \leftarrow \mathbf{W}^{[k]T} \cdot \nabla L(\mathbf{Z}^{[k]})$

end

for $k \leftarrow 1$ **to** L **do**

$\mathbf{W}^{[k]} \leftarrow \mathbf{W}^{[k]} - \gamma \nabla L(\mathbf{W}^{[k]})$ *[Update parameters]*

end

end

Batch gradient descent

If the training set is too large, it takes a lot of time to process the whole training set for each step of gradient descent.

Mini-batch gradient descent

If the training set is too large, it takes a lot of time to process the whole training set for each step of gradient descent.

→ Divide the training set into multiple mini-batches of **reasonable** sizes.

$$\begin{aligned}\mathbf{X} &= [\mathbf{x}^{(1)} \quad \mathbf{x}^{(2)} \dots \mathbf{x}^{(s)} \quad | \quad \mathbf{x}^{(s+1)} \dots \mathbf{x}^{(2s)} \quad | \quad \mathbf{x}^{(2s+1)} \dots \quad | \quad \dots \mathbf{x}^{(N)}] \\ &= [\mathbf{X}_{\{1\}} \quad | \quad \mathbf{X}_{\{2\}} \quad | \quad \dots \quad | \quad \mathbf{X}_{\{N/s\}}] \\ \mathbf{Y} &= [y^1 \quad y^2 \dots y^s \quad | \quad y^{s+1} \dots y^{2s} \quad | \quad y^{2s+1} \dots \quad | \quad \dots y^N] \\ &= [\mathbf{Y}_{\{1\}} \quad | \quad \mathbf{Y}_{\{2\}} \quad | \quad \dots \quad | \quad \mathbf{Y}_{\{N/s\}}]\end{aligned}$$

At each iteration t , perform gradient descent using one mini-batch of training examples $(\mathbf{X}_{\{t\}}; \mathbf{Y}_{\{t\}})$.

Mini-batch gradient descent

- Mini-batch size $s = N \longrightarrow$ Normal gradient descent.
- Mini-batch size $s = 1 \longrightarrow$ Stochastic gradient descent.
- Mini-batch sizes are normally set as powers of 2, e.g.,
 $s = 64, 128, 256, 512$

Problem with gradient descent




Figure: Gradient descent ²

²<https://engmrk.com/gradient-descent-with-momentum/>.

Problem with gradient descent





Figure: Gradient descent ³


³<https://engmrk.com/gradient-descent-with-momentum/>.

Exponentially weighted moving average

Example: How to model/forecast the stock/share price of a certain company?



Exponentially weighted moving average



$$m_{100} = \beta m_{99} + (1 - \beta) price_{100}$$

$$m_{99} = \beta m_{98} + (1 - \beta) price_{99}$$


...

$$m_t = \beta m_{t-1} + (1 - \beta) price_t$$

...

$$m_1 = \beta m_0 + (1 - \beta) price_1$$

Exponentially weighted moving average



$$\beta = 0.9$$

$$m_{100} = 0.9m_{99} + (1 - 0.9)price_{100}$$

$$= 0.1price_{100} + 0.9m_{99}$$

$$= 0.1price_{100} + 0.9(0.1price_{99} + 0.9m_{98})$$


$$= 0.1price_{100} + 0.9 \cdot 0.1price_{99} + (0.9)^2 m_{98}$$

$$= 0.1price_{100} + 0.09price_{99} + (0.9)^2 (0.1price_{98} + 0.9m_{97})$$

$$= 0.1price_{100} + 0.09price_{99} + 0.081price_{98} + (0.9)^3 m_{97}$$

$$= \dots$$

Exponentially weighted moving average



Actually, without the below correction, we will get the orange line instead of the red line. Why?

$$m_t = \beta m_{t-1} + (1 - \beta) \text{price}_t$$
$$m_t^{\text{corrected}} = \frac{m_t}{1 - \beta^t}$$

Gradient descent with momentum

Update weights at iteration t :

$$\mathbf{g}_{\{t\}} \leftarrow \nabla L(\mathbf{W}_{\{t-1\}})$$

$$\mathbf{m}_{\{t\}} \leftarrow \beta_1 \mathbf{m}_{\{t-1\}} + (1 - \beta_1) \mathbf{g}_{\{t\}}$$

$$\mathbf{W}_{\{t\}} \leftarrow \mathbf{W}_{\{t-1\}} - \gamma \mathbf{m}_{\{t\}}$$

Initialization: $\mathbf{m}_{\{0\}} = \mathbf{0}$

Hyperparameters: $\beta_1 = 0.9$; learning rate γ

Gradient descent with momentum




Figure: Gradient descent with momentum ⁴

⁴<https://engmrk.com/gradient-descent-with-momentum/>.

Problem with gradient descent




Figure: Gradient descent ⁵

⁵https://github.com/hengluchang/visualizing_momentum.

Gradient descent with momentum




Figure: Gradient descent with momentum $\beta_1 = 0.8$ ⁶

⁶https://github.com/hengluchang/visualizing_momentum.

Update weights at iteration t :

$$\begin{aligned}\mathbf{g}_{\{t\}} &\leftarrow \nabla L(\mathbf{W}_{\{t-1\}}) \\ \mathbf{s}_{\{t\}} &\leftarrow \beta_2 \mathbf{s}_{\{t-1\}} + (1 - \beta_2) \mathbf{g}_t^2 \\ \mathbf{W}_{\{t\}} &\leftarrow \mathbf{W}_{\{t-1\}} - \gamma \frac{\mathbf{g}_{\{t\}}}{\sqrt{\mathbf{s}_{\{t\}} + \epsilon}}\end{aligned}$$

Initialization: $\mathbf{s}_{\{0\}} = \mathbf{0}$

Hyperparameters: $\beta_2 = 0.9$; $\epsilon = 10^{-8}$; learning rate γ

Adam - Adaptive moment estimation

Update weights at iteration t :

$$\mathbf{g}_{\{t\}} \leftarrow \nabla L(\mathbf{W}_{\{t-1\}})$$

$$\mathbf{m}_{\{t\}} \leftarrow \beta_1 \mathbf{m}_{\{t-1\}} + (1 - \beta_1) \mathbf{g}_{\{t\}}$$

$$\mathbf{s}_{\{t\}} \leftarrow \beta_2 \mathbf{s}_{\{t-1\}} + (1 - \beta_2) \mathbf{g}_{\{t\}}^2$$

$$\mathbf{m}_{\{t\}}^{\text{corrected}} \leftarrow \frac{\mathbf{m}_{\{t\}}}{1 - \beta_1^t}$$

$$\mathbf{s}_{\{t\}}^{\text{corrected}} \leftarrow \frac{\mathbf{s}_{\{t\}}}{1 - \beta_2^t}$$

$$\mathbf{W}_{\{t\}} \leftarrow \mathbf{W}_{\{t-1\}} - \gamma \frac{\mathbf{m}_{\{t\}}^{\text{corrected}}}{\sqrt{\mathbf{s}_{\{t\}}^{\text{corrected}} + \epsilon}}$$

Initialization: $\mathbf{m}_{\{0\}} = \mathbf{s}_{\{0\}} = \mathbf{0}$

Hyperparameters: $\beta_1 = 0.9$; $\beta_2 = 0.999$; $\epsilon = 10^{-8}$; learning rate γ

- ① Visualizing Gradient Descent with Momentum in Python:
https://github.com/hengluchang/visualizing_momentum.
- ② Coursera - Improving Deep Neural Networks: Hyperparameter tuning, Regularization and Optimization:
<https://www.coursera.org/learn/deep-neural-network/>
- ③ Coursera - Neural Networks and Deep Learning:
<https://www.coursera.org/learn/neural-networks-deep-learning/>